

Actividad2, Técnicas Multivariantes

Jorge A. Balsells Orellana

May 17, 2021

En esta actividad vas a profundizar en las distintas técnicas que se pueden aplicar para abordar un problema de regresión. Además, profundizarás en tus conocimientos sobre las librerías statsmodels y scikit-learn de Python.

0.1 Importando librerías

```
[1]: from sklearn.model_selection import train_test_split
from sklearn.datasets import make_regression
from sklearn import linear_model
import matplotlib.pyplot as plt
import statsmodels.api as sm
import sklearn as sk
import numpy as np
import pandas as pd
```

1 Creando dataset

El primer paso consiste en crear un conjunto de datos ficticio. Para garantizar que cada alumno obtiene uno distinto se va a emplear el documento de identidad de cada uno para crear el conjunto de datos. Para que sean comparables entre todos, si el número de identidad tiene menos de 8 cifras replicaremos las primeras hasta obtener exactamente 8. Además, para evitar los dígitos 0 y 1, si alguna de las cifras es menor que 2 la sustituiremos por ese número. Aplicando estos cambios tendremos el número del documento de identidad preparado para la resolución de la actividad. Numero de identificación personal: 166315389-0101, equivale a 26632538

```
[2]: data = sk.datasets.make_regression(
    n_samples=200+10*2,
    n_features=10+6+6,
    n_informative=10+6,
    n_targets=1,
    noise=10*3,
    shuffle=False,
    bias=2,
    random_state=None, effective_rank=None,
    tail_strength=0.5, coef=False
)
```

2 Agregando títulos

```
[3]: data_fr = pd.DataFrame(data[0])
data_fr.columns=['c0','c1','c2','c3','c4','c5','c6','c7','c8','c9','c10','c11','c12','c13','c14','c15','c16','c17','c18','c19','c20','c21']
print(data_fr)
```

	c0	c1	c2	c3	c4	c5	c6 \
0	-1.200619	0.363886	0.527277	1.776005	1.562085	-0.998717	-1.445892
1	-0.370290	-0.721044	-0.004600	1.988354	-1.187087	1.929338	1.106362
2	-0.083060	-0.848515	-0.666145	2.155664	1.217626	0.442125	-0.066560
3	0.547946	0.137935	0.803414	1.647611	-0.406032	-2.838325	0.998711
4	-0.478050	0.939166	1.590225	-2.444970	-1.545710	0.086851	-0.821011
..
215	0.076349	-0.294829	0.649279	-1.841271	0.515866	-0.785015	0.919299
216	-1.447078	0.530184	0.876954	-0.646612	0.652381	-0.302676	-1.117376
217	-0.197006	-2.568405	0.944425	1.424370	1.913104	0.213528	0.841183
218	0.148215	1.986257	1.706012	1.239836	-0.204384	-1.498575	1.656196
219	0.101630	1.352317	0.540504	1.114844	-0.140460	-1.433073	1.174183

	c7	c8	c9	...	c12	c13	c14 \
0	1.007802	-0.210287	-0.315635	...	-0.846384	0.751240	-0.831247
1	-1.119966	1.421816	-0.070631	...	-0.386086	0.305581	1.619579
2	-0.238667	0.858434	-0.115024	...	0.949134	-1.379160	1.224498
3	-1.523243	-0.242737	-1.103062	...	0.449596	-2.452147	-1.325093
4	0.187369	-0.064385	0.240592	...	0.141086	-0.828289	1.132770
..
215	-0.460281	1.771298	-0.550230	...	-0.181780	-1.344381	-0.942128
216	-0.952387	0.646239	-0.520901	...	-0.118848	-0.840821	0.817320
217	1.971822	1.805175	-1.793488	...	0.081922	-0.748347	-1.119656
218	1.336042	-0.428932	0.776098	...	2.081237	-0.391987	-0.439679
219	1.167033	-0.296465	-1.497535	...	0.890979	-0.105300	1.564012

	c15	c16	c17	c18	c19	c20	c21
0	-2.234769	-1.388944	0.341197	1.684392	-0.810816	0.172313	-0.943393
1	0.406285	-1.096114	-0.904892	-1.189931	-0.551215	0.579500	-2.862278
2	0.048189	-1.182454	-1.991184	-0.456665	-0.371789	-0.618690	-0.531200
3	-0.034923	0.227892	0.140545	-0.679035	-1.118085	0.325487	-1.751562
4	-2.100005	-0.793645	1.414209	0.430947	-1.358468	1.061865	0.168035
..
215	-0.763008	0.103196	0.059206	0.813484	0.177341	0.376273	0.912228
216	0.154747	0.249033	0.124586	0.790569	1.317087	1.068314	-0.180360
217	-1.136287	-1.632205	-0.185916	0.315124	0.350710	0.078135	-1.662309
218	-0.825360	-0.954144	0.378729	1.972238	-0.975027	-0.040878	-0.707584
219	0.717811	-0.886581	0.015919	0.702115	0.673015	-1.029820	-0.050169

[220 rows x 22 columns]

3 Dividiendo Dataframe a través de un Split

Divide el conjunto de datos en 200 observaciones para el entrenamiento y el resto para realizar la validación de los distintos métodos de regresión aplicados.

```
[4]: train_set, test_set, y_train, y_test = train_test_split(data_fr, data[1],  
    →test_size=0.09, random_state=42)  
  
print(f"tamaño entrenamiento features {train_set.shape}")  
print(f"tamaño test features {test_set.shape}")  
print(f"tamaño entrenamiento respuesta {y_train.shape}")  
print(f"tamaño test respuesta {y_test.shape}")
```

```
tamaño entrenamiento features (200, 22)  
tamaño test features (20, 22)  
tamaño entrenamiento respuesta (200,)  
tamaño test respuesta (20,)
```

Describe tu conjunto de datos (transformalo en un data.frame, aplica los métodos .info(), .describe()) y obtén el histograma de todas las variables (predictoras y la variable respuesta).

```
[5]: train_set.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 200 entries, 189 to 102  
Data columns (total 22 columns):  
#   Column  Non-Null Count  Dtype  
---  ---  
0   c0      200 non-null    float64  
1   c1      200 non-null    float64  
2   c2      200 non-null    float64  
3   c3      200 non-null    float64  
4   c4      200 non-null    float64  
5   c5      200 non-null    float64  
6   c6      200 non-null    float64  
7   c7      200 non-null    float64  
8   c8      200 non-null    float64  
9   c9      200 non-null    float64  
10  c10     200 non-null    float64  
11  c11     200 non-null    float64  
12  c12     200 non-null    float64  
13  c13     200 non-null    float64  
14  c14     200 non-null    float64  
15  c15     200 non-null    float64  
16  c16     200 non-null    float64  
17  c17     200 non-null    float64
```

```

18  c18      200 non-null    float64
19  c19      200 non-null    float64
20  c20      200 non-null    float64
21  c21      200 non-null    float64
dtypes: float64(22)
memory usage: 35.9 KB

```

```
[6]: train_set.describe()
```

```

[6]:
count      c0      c1      c2      c3      c4      c5  \
count  200.000000  200.000000  200.000000  200.000000  200.000000  200.000000
mean    0.061775  0.068480  -0.029100  -0.079289  0.024882  0.058193
std     0.951104  0.859848  0.935541  1.045438  1.029154  0.990542
min    -2.257323  -2.568405  -2.482018  -2.660958  -3.065924  -2.838325
25%    -0.533241  -0.546414  -0.606245  -0.740814  -0.688340  -0.597589
50%     0.102723  0.133160  -0.073856  -0.105574  0.089689  0.072999
75%     0.744515  0.681227  0.627146  0.691111  0.691641  0.776711
max     2.827125  2.731737  2.353466  2.578257  2.533175  2.936126

count      c6      c7      c8      c9  ...      c12  \
count  200.000000  200.000000  200.000000  200.000000  ...  200.000000
mean   -0.061592  0.020946  0.024503  -0.093435  ...  -0.072567
std     0.898575  1.046160  0.983612  0.951497  ...  1.004430
min    -2.377142  -2.372431  -2.582735  -2.757625  ...  -2.571439
25%    -0.645426  -0.744354  -0.657205  -0.675275  ...  -0.668238
50%    -0.062011  -0.085418  0.019527  -0.079579  ...  -0.108140
75%     0.584005  0.664773  0.690342  0.540392  ...  0.602127
max     2.180832  3.437404  2.713194  2.434474  ...  3.605012

count      c13      c14      c15      c16      c17      c18  \
count  200.000000  200.000000  200.000000  200.000000  200.000000  200.000000
mean    0.018591  0.064096  -0.042001  0.051858  -0.066966  -0.010057
std     1.008430  0.987427  0.961506  1.068868  0.910952  1.071814
min    -2.511117  -2.603918  -2.434906  -3.314557  -2.959647  -2.939178
25%    -0.699544  -0.597825  -0.652038  -0.604215  -0.633148  -0.712945
50%     0.082356  0.013870  -0.005546  0.034842  -0.064684  0.023465
75%     0.686061  0.741414  0.602292  0.708587  0.621575  0.614373
max     3.105076  2.967300  2.438551  2.630969  2.541617  2.675915

count      c19      c20      c21
count  200.000000  200.000000  200.000000
mean   -0.048863  0.036259  -0.090818
std     0.893632  0.995174  0.973091
min    -2.101485  -2.593781  -2.862278
25%    -0.590866  -0.631175  -0.696132
50%    -0.114663  0.061491  -0.049538
75%     0.520805  0.803799  0.533666

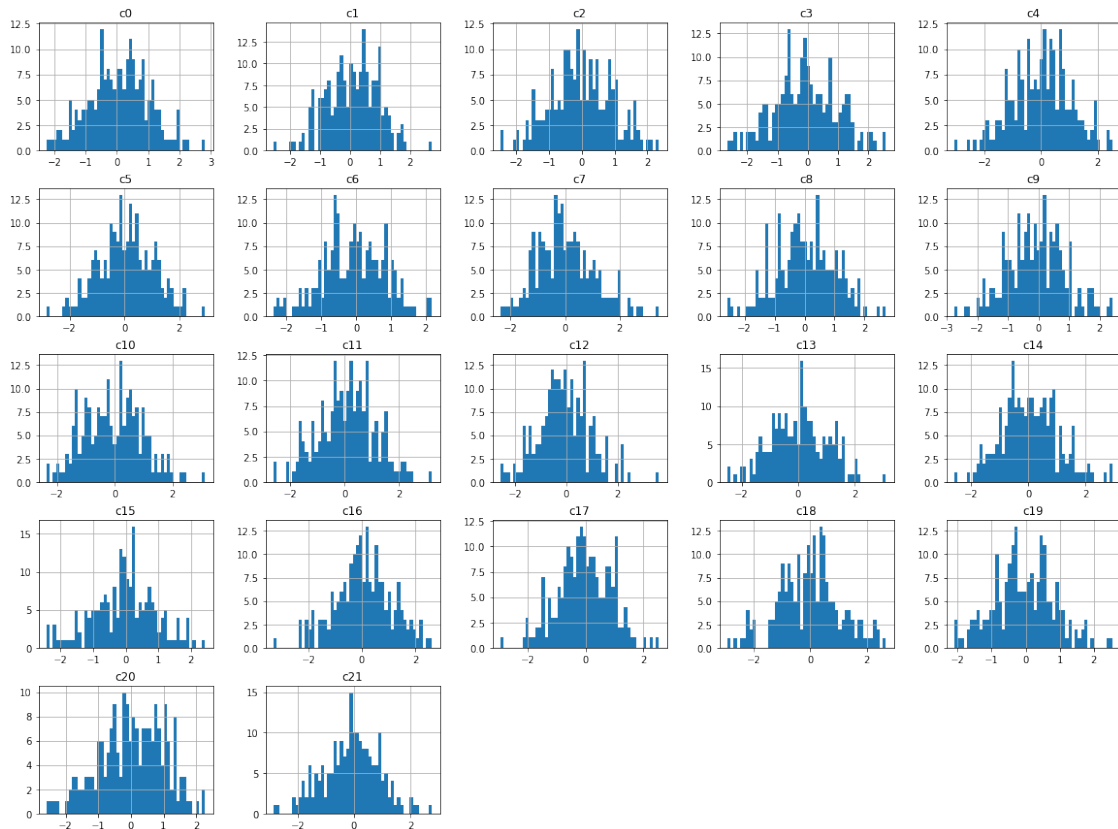
```

```
max      2.559372    2.276325    2.769567
```

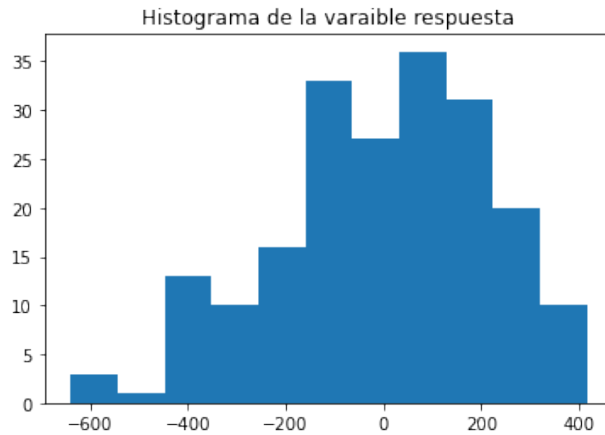
```
[8 rows x 22 columns]
```

4 Histogramas de features

```
[7]: train_set.hist(bins=50, figsize=(20,15))  
plt.show()
```



```
[8]: plt.hist(y_train, bins='auto')  
plt.title("Histograma de la variable respuesta")  
plt.show()
```



5 Regresión lineal múltiple

Obtén un modelo de regresión lineal múltiple. ¿Son todas las variables predictoras significativas? Utiliza la librería statsmodels. **No son todas las variables predictoras significativas, ya que algunas no cumplen con $P > |t|$.**

```
[9]: train_data = train_set.copy()
train_data['y'] = y_train
mat_corr = train_data.corr()
mat_corr['y']
```

```
[9]: c0      0.364644
c1      0.250325
c2     -0.107469
c3      0.424363
c4     -0.037278
c5      0.351267
c6      0.322329
c7      0.231789
c8     -0.018025
c9      0.251766
c10     0.130534
c11     0.168569
c12     0.118699
c13     0.161102
c14     0.437383
c15     0.233611
c16     0.062251
c17    -0.133591
c18    -0.017822
c19     0.024090
```

```

c20    -0.070258
c21     0.031895
y       1.000000
Name: y, dtype: float64

```

6 Entrenamiento modelo 1

```

[11]: train_set = sm.add_constant(train_set, prepend=True)
      model1 = sm.OLS(endog=y_train, exog=train_set)
      model1 = model1.fit()

```

```

[12]: print(model1.summary())

```

```

                    OLS Regression Results
=====
Dep. Variable:            y    R-squared:            0.982
Model:                  OLS    Adj. R-squared:        0.979
Method:             Least Squares    F-statistic:            430.1
Date:                Mon, 17 May 2021    Prob (F-statistic):       2.78e-141
Time:                  01:34:36    Log-Likelihood:          -959.93
No. Observations:        200    AIC:                   1966.
Df Residuals:            177    BIC:                   2042.
Df Model:                 22
Covariance Type:            nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-0.2033	2.326	-0.087	0.930	-4.794	4.388
c0	71.3972	2.483	28.758	0.000	66.498	76.297
c1	65.7305	2.717	24.189	0.000	60.368	71.093
c2	2.6603	2.580	1.031	0.304	-2.431	7.751
c3	96.9798	2.290	42.354	0.000	92.461	101.499
c4	17.8190	2.313	7.705	0.000	13.255	22.383
c5	55.8696	2.350	23.774	0.000	51.232	60.507
c6	65.0975	2.617	24.878	0.000	59.934	70.261
c7	42.9200	2.204	19.472	0.000	38.570	47.270
c8	7.8005	2.340	3.334	0.001	3.183	12.418
c9	62.1009	2.541	24.439	0.000	57.086	67.116
c10	22.1543	2.280	9.719	0.000	17.656	26.653
c11	18.4234	2.227	8.272	0.000	14.028	22.819
c12	31.9253	2.329	13.706	0.000	27.329	36.522
c13	39.7516	2.318	17.148	0.000	35.177	44.326
c14	98.2821	2.323	42.301	0.000	93.697	102.867
c15	45.2187	2.473	18.288	0.000	40.339	50.098
c16	2.7655	2.168	1.276	0.204	-1.513	7.044
c17	2.6992	2.632	1.025	0.307	-2.496	7.894

c18	2.1383	2.263	0.945	0.346	-2.328	6.604
c19	0.7532	2.523	0.299	0.766	-4.225	5.731
c20	-3.2485	2.373	-1.369	0.173	-7.931	1.434
c21	-0.7817	2.521	-0.310	0.757	-5.756	4.193

Omnibus:	6.271	Durbin-Watson:	1.991
Prob(Omnibus):	0.043	Jarque-Bera (JB):	9.922
Skew:	0.045	Prob(JB):	0.00701
Kurtosis:	4.087	Cond. No.	2.01

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

7 Aplicando Step-Wise

Realiza una selección de variables mediante un algoritmo de tipo step-wise, donde en cada paso elimines la variable predictora menos significativa atendiendo a su p.valor hasta que todas las variables del modelo sean significativas ($p.\text{valor} < 0.05$).

En este caso, el valor de $P > |t|$ debe cumplir, siendo P menor a $1.35 * e^{-142} \dots$ A continuación se eliminan todas las columnas del train-set que correspondan a valores mayores de P . En este caso, corresponde inicialmente a la columna **C19, C21, C18, C17, C20, C2 y C16**.

```
[13]: train_set.drop('c19', axis='columns', inplace=True)
train_set = sm.add_constant(train_set, prepend=True)
model2 = sm.OLS(endog=y_train, exog=train_set)
model2 = model2.fit()
```

```
[14]: print(model2.summary())
```

```

OLS Regression Results
=====
Dep. Variable:          y      R-squared:          0.982
Model:                  OLS    Adj. R-squared:      0.979
Method:                 Least Squares    F-statistic:      452.9
Date:                  Mon, 17 May 2021    Prob (F-statistic): 1.35e-142
Time:                  01:34:53    Log-Likelihood:    -959.98
No. Observations:      200    AIC:              1964.
Df Residuals:          178    BIC:              2037.
Df Model:              21
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-0.2604	2.313	-0.113	0.910	-4.824	4.303
c0	71.4167	2.476	28.849	0.000	66.532	76.302

c1	65.7940	2.702	24.349	0.000	60.462	71.126
c2	2.6727	2.573	1.039	0.300	-2.404	7.750
c3	96.9954	2.283	42.480	0.000	92.490	101.501
c4	17.8809	2.297	7.783	0.000	13.347	22.414
c5	55.8882	2.343	23.851	0.000	51.264	60.512
c6	65.0868	2.610	24.940	0.000	59.937	70.237
c7	42.8624	2.190	19.571	0.000	38.540	47.184
c8	7.8136	2.333	3.349	0.001	3.209	12.418
c9	62.0935	2.534	24.500	0.000	57.092	67.095
c10	22.1471	2.274	9.741	0.000	17.660	26.634
c11	18.4483	2.220	8.310	0.000	14.067	22.829
c12	31.8587	2.313	13.776	0.000	27.295	36.422
c13	39.7645	2.312	17.200	0.000	35.202	44.327
c14	98.2918	2.317	42.418	0.000	93.719	102.865
c15	45.2141	2.466	18.333	0.000	40.347	50.081
c16	2.8038	2.158	1.299	0.196	-1.456	7.063
c17	2.6858	2.625	1.023	0.308	-2.495	7.867
c18	2.1186	2.256	0.939	0.349	-2.334	6.571
c20	-3.2184	2.364	-1.361	0.175	-7.884	1.447
c21	-0.7839	2.514	-0.312	0.756	-5.746	4.178

```
=====
Omnibus:                6.547    Durbin-Watson:                1.988
Prob(Omnibus):          0.038    Jarque-Bera (JB):         10.683
Skew:                   0.033    Prob(JB):                 0.00479
Kurtosis:               4.130    Cond. No.                 1.99
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[15]: train_set.drop('c21', axis='columns', inplace=True)
      train_set = sm.add_constant(train_set, prepend=True)
      model2 = sm.OLS(endog=y_train, exog=train_set)
      model2 = model2.fit()
```

```
[16]: print(model2.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          y    R-squared:                0.982
Model:                  OLS    Adj. R-squared:         0.980
Method:                 Least Squares    F-statistic:          478.0
Date:                   Mon, 17 May 2021    Prob (F-statistic):    6.39e-144
Time:                   01:35:05    Log-Likelihood:       -960.03
No. Observations:      200    AIC:                  1962.
Df Residuals:          179    BIC:                  2031.
Df Model:              20
```

Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	-0.1809	2.293	-0.079	0.937	-4.705	4.343
c0	71.2484	2.410	29.566	0.000	66.493	76.004
c1	65.8150	2.694	24.426	0.000	60.498	71.132
c2	2.7500	2.554	1.077	0.283	-2.290	7.790
c3	97.1219	2.241	43.332	0.000	92.699	101.545
c4	17.9285	2.286	7.841	0.000	13.417	22.440
c5	55.9942	2.313	24.213	0.000	51.431	60.558
c6	65.1156	2.602	25.030	0.000	59.982	70.249
c7	42.8445	2.184	19.619	0.000	38.535	47.154
c8	7.8812	2.317	3.401	0.001	3.308	12.454
c9	61.9033	2.454	25.228	0.000	57.061	66.745
c10	22.2037	2.261	9.822	0.000	17.743	26.665
c11	18.4284	2.214	8.325	0.000	14.060	22.796
c12	31.8425	2.306	13.807	0.000	27.292	36.393
c13	39.7847	2.305	17.259	0.000	35.236	44.333
c14	98.2064	2.295	42.789	0.000	93.677	102.735
c15	45.3165	2.438	18.587	0.000	40.505	50.128
c16	2.7531	2.147	1.282	0.201	-1.483	6.990
c17	2.6668	2.618	1.019	0.310	-2.499	7.833
c18	2.0328	2.234	0.910	0.364	-2.375	6.441
c20	-3.1425	2.346	-1.340	0.182	-7.772	1.487
Omnibus:		6.356	Durbin-Watson:			1.988
Prob(Omnibus):		0.042	Jarque-Bera (JB):			10.201
Skew:		0.030	Prob(JB):			0.00609
Kurtosis:		4.105	Cond. No.			1.88

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[17]: train_set.drop('c18', axis='columns', inplace=True)
train_set = sm.add_constant(train_set, prepend=True)
model2 = sm.OLS(endog=y_train, exog=train_set)
model2 = model2.fit()
```

```
[18]: print(model2.summary())
```

OLS Regression Results

Dep. Variable:	y	R-squared:	0.982
Model:	OLS	Adj. R-squared:	0.980
Method:	Least Squares	F-statistic:	503.6

Date: Mon, 17 May 2021 Prob (F-statistic): 4.24e-145
Time: 01:35:19 Log-Likelihood: -960.50
No. Observations: 200 AIC: 1961.
Df Residuals: 180 BIC: 2027.
Df Model: 19
Covariance Type: nonrobust

	coef	std err	t	P> t	[0.025	0.975]
const	-0.3254	2.286	-0.142	0.887	-4.836	4.186
c0	71.3017	2.408	29.611	0.000	66.550	76.053
c1	65.8078	2.693	24.435	0.000	60.494	71.122
c2	3.0360	2.534	1.198	0.232	-1.964	8.036
c3	97.3354	2.228	43.689	0.000	92.939	101.732
c4	17.9433	2.285	7.852	0.000	13.434	22.453
c5	56.0758	2.310	24.278	0.000	51.518	60.633
c6	64.8721	2.586	25.081	0.000	59.768	69.976
c7	42.7109	2.178	19.611	0.000	38.413	47.008
c8	8.0906	2.305	3.510	0.001	3.542	12.639
c9	61.4269	2.396	25.636	0.000	56.699	66.155
c10	21.9495	2.242	9.789	0.000	17.525	26.374
c11	18.5601	2.208	8.407	0.000	14.204	22.916
c12	31.4831	2.271	13.863	0.000	27.002	35.964
c13	40.0765	2.282	17.565	0.000	35.574	44.579
c14	98.2379	2.294	42.828	0.000	93.712	102.764
c15	45.5399	2.425	18.783	0.000	40.756	50.324
c16	2.8142	2.145	1.312	0.191	-1.418	7.047
c17	2.5329	2.613	0.969	0.334	-2.622	7.688
c20	-2.7757	2.310	-1.202	0.231	-7.334	1.782
Omnibus:	6.022	Durbin-Watson:	1.989			
Prob(Omnibus):	0.049	Jarque-Bera (JB):	9.415			
Skew:	0.007	Prob(JB):	0.00903			
Kurtosis:	4.063	Cond. No.	1.78			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[19]: train_set.drop('c17', axis='columns', inplace=True)
train_set = sm.add_constant(train_set, prepend=True)
model2 = sm.OLS(endog=y_train, exog=train_set)
model2 = model2.fit()
```

```
[20]: print(model2.summary())
```

OLS Regression Results

```

=====
Dep. Variable:          y      R-squared:          0.981
Model:                  OLS    Adj. R-squared:       0.980
Method:                 Least Squares  F-statistic:       531.7
Date:                  Mon, 17 May 2021  Prob (F-statistic): 2.89e-146
Time:                  01:35:32  Log-Likelihood:    -961.02
No. Observations:      200     AIC:              1960.
Df Residuals:          181     BIC:              2023.
Df Model:              18
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-0.5507	2.274	-0.242	0.809	-5.037	3.936
c0	71.0005	2.387	29.739	0.000	66.290	75.711
c1	65.4634	2.669	24.526	0.000	60.197	70.730
c2	2.9783	2.533	1.176	0.241	-2.019	7.975
c3	96.8253	2.165	44.732	0.000	92.554	101.096
c4	18.0813	2.280	7.929	0.000	13.582	22.581
c5	56.1468	2.308	24.325	0.000	51.592	60.701
c6	64.8758	2.586	25.087	0.000	59.773	69.978
c7	42.9235	2.166	19.813	0.000	38.649	47.198
c8	7.9689	2.301	3.463	0.001	3.428	12.509
c9	61.4172	2.396	25.637	0.000	56.690	66.144
c10	21.9190	2.242	9.778	0.000	17.496	26.342
c11	18.8089	2.192	8.579	0.000	14.483	23.135
c12	31.2239	2.255	13.847	0.000	26.775	35.673
c13	40.1395	2.280	17.603	0.000	35.640	44.639
c14	98.2658	2.293	42.851	0.000	93.741	102.791
c15	45.4753	2.423	18.767	0.000	40.694	50.257
c16	2.6426	2.137	1.237	0.218	-1.574	6.860
c20	-2.6741	2.307	-1.159	0.248	-7.226	1.878

```

=====
Omnibus:                5.772    Durbin-Watson:          1.997
Prob(Omnibus):          0.056    Jarque-Bera (JB):        8.784
Skew:                   0.024    Prob(JB):                0.0124
Kurtosis:               4.026    Cond. No.:               1.74
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

[21]: train_set.drop('c20', axis='columns', inplace=True)
      train_set = sm.add_constant(train_set, prepend=True)
      model2 = sm.OLS(endog=y_train, exog=train_set)
      model2 = model2.fit()

```

```
[22]: print(model2.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  y      R-squared:                  0.981
Model:                        OLS      Adj. R-squared:              0.980
Method:                    Least Squares  F-statistic:                  561.8
Date:                Mon, 17 May 2021  Prob (F-statistic):          2.34e-147
Time:                  01:35:44      Log-Likelihood:              -961.76
No. Observations:          200      AIC:                          1960.
Df Residuals:              182      BIC:                          2019.
Df Model:                   17
Covariance Type:            nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-0.6684	2.274	-0.294	0.769	-5.155	3.818
c0	70.9421	2.389	29.694	0.000	66.228	75.656
c1	65.6803	2.665	24.645	0.000	60.422	70.939
c2	3.1229	2.532	1.233	0.219	-1.873	8.119
c3	97.1546	2.148	45.233	0.000	92.917	101.392
c4	18.5179	2.251	8.225	0.000	14.076	22.960
c5	56.0162	2.308	24.275	0.000	51.463	60.569
c6	64.7596	2.587	25.037	0.000	59.656	69.863
c7	42.8632	2.168	19.772	0.000	38.586	47.141
c8	8.0017	2.303	3.474	0.001	3.457	12.546
c9	61.2205	2.392	25.595	0.000	56.501	65.940
c10	22.0990	2.238	9.873	0.000	17.683	26.515
c11	18.8339	2.194	8.583	0.000	14.504	23.164
c12	31.1036	2.255	13.795	0.000	26.655	35.552
c13	40.0868	2.282	17.566	0.000	35.584	44.589
c14	98.4499	2.290	42.994	0.000	93.932	102.968
c15	45.7302	2.415	18.932	0.000	40.964	50.496
c16	2.8738	2.130	1.349	0.179	-1.329	7.076

```

=====
Omnibus:                      6.004      Durbin-Watson:              1.997
Prob(Omnibus):                0.050      Jarque-Bera (JB):           9.310
Skew:                         0.034      Prob(JB):                   0.00952
Kurtosis:                     4.055      Cond. No.                    1.68
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[23]: train_set.drop('c2', axis='columns', inplace=True)
      train_set = sm.add_constant(train_set, prepend=True)
```

```
model2 = sm.OLS(endog=y_train, exog=train_set)
model2 = model2.fit()
```

```
[24]: print(model2.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  y      R-squared:                0.981
Model:                        OLS      Adj. R-squared:           0.979
Method:                    Least Squares  F-statistic:             595.1
Date:                Mon, 17 May 2021  Prob (F-statistic):       2.01e-148
Time:                        01:35:58  Log-Likelihood:          -962.59
No. Observations:                200   AIC:                     1959.
Df Residuals:                    183   BIC:                     2015.
Df Model:                        16
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-0.6630	2.277	-0.291	0.771	-5.155	3.829
c0	70.6606	2.382	29.669	0.000	65.962	75.360
c1	65.4699	2.663	24.581	0.000	60.215	70.725
c3	97.0777	2.150	45.152	0.000	92.836	101.320
c4	18.0597	2.224	8.122	0.000	13.672	22.447
c5	55.5940	2.285	24.327	0.000	51.085	60.103
c6	64.7329	2.590	24.992	0.000	59.623	69.843
c7	42.8066	2.170	19.723	0.000	38.524	47.089
c8	7.7933	2.300	3.388	0.001	3.255	12.332
c9	61.5895	2.376	25.916	0.000	56.901	66.278
c10	22.2372	2.239	9.933	0.000	17.820	26.654
c11	19.2248	2.174	8.841	0.000	14.935	23.515
c12	31.1868	2.257	13.819	0.000	26.734	35.640
c13	39.5467	2.243	17.633	0.000	35.122	43.972
c14	98.3032	2.290	42.927	0.000	93.785	102.821
c15	45.5222	2.413	18.865	0.000	40.761	50.283
c16	2.7093	2.129	1.273	0.205	-1.491	6.909

```

=====
Omnibus:                    5.220   Durbin-Watson:           2.021
Prob(Omnibus):              0.074   Jarque-Bera (JB):        7.417
Skew:                      0.052   Prob(JB):                0.0245
Kurtosis:                  3.938   Cond. No.                 1.61
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[26]: train_set.drop('c16', axis='columns', inplace=True)
train_set = sm.add_constant(train_set, prepend=True)
model2 = sm.OLS(endog=y_train, exog=train_set)
model2 = model2.fit()
```

```
[27]: print(model2.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  y      R-squared:                0.981
Model:                        OLS      Adj. R-squared:           0.979
Method:                    Least Squares      F-statistic:           632.6
Date:                Mon, 17 May 2021      Prob (F-statistic):       1.75e-149
Time:                        01:36:18      Log-Likelihood:          -963.47
No. Observations:                200      AIC:                    1959.
Df Residuals:                    184      BIC:                    2012.
Df Model:                        15
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-0.4821	2.276	-0.212	0.832	-4.973	4.009
c0	70.6325	2.386	29.609	0.000	65.926	75.339
c1	65.7698	2.657	24.749	0.000	60.527	71.013
c3	97.0603	2.154	45.069	0.000	92.811	101.309
c4	18.0374	2.227	8.098	0.000	13.643	22.432
c5	55.5705	2.289	24.276	0.000	51.054	60.087
c6	64.6963	2.594	24.938	0.000	59.578	69.815
c7	43.0078	2.168	19.835	0.000	38.730	47.286
c8	7.7346	2.304	3.358	0.001	3.190	12.279
c9	61.6922	2.379	25.931	0.000	56.998	66.386
c10	22.0828	2.239	9.862	0.000	17.665	26.501
c11	19.0877	2.175	8.774	0.000	14.796	23.380
c12	31.6224	2.234	14.152	0.000	27.214	36.031
c13	39.3717	2.242	17.558	0.000	34.948	43.796
c14	98.2735	2.294	42.844	0.000	93.748	102.799
c15	45.8517	2.403	19.080	0.000	41.110	50.593

```

=====
Omnibus:                    5.633      Durbin-Watson:           2.036
Prob(Omnibus):              0.060      Jarque-Bera (JB):         8.378
Skew:                      0.046      Prob(JB):                 0.0152
Kurtosis:                  3.999      Cond. No.                  1.60
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

8 Regresión mediante la red elástica

Realiza una regresión mediante la red elástica. Prueba distintos valores de α y obtén el valor óptimo de α y de λ mediante validación cruzada.

```
[28]: from sklearn.model_selection import RepeatedKFold
      from sklearn.model_selection import GridSearchCV
      from sklearn.metrics import mean_squared_error
      from sklearn.linear_model import ElasticNet
```

```
[29]: parameters = {
      'alpha': np.arange(0.005, 0.9, 0.01),
      'l1_ratio': np.arange(0.1, 0.5, 0.01)
      }
      print(parameters)
```

```
{'alpha': array([0.005, 0.015, 0.025, 0.035, 0.045, 0.055, 0.065, 0.075, 0.085,
                0.095, 0.105, 0.115, 0.125, 0.135, 0.145, 0.155, 0.165, 0.175,
                0.185, 0.195, 0.205, 0.215, 0.225, 0.235, 0.245, 0.255, 0.265,
                0.275, 0.285, 0.295, 0.305, 0.315, 0.325, 0.335, 0.345, 0.355,
                0.365, 0.375, 0.385, 0.395, 0.405, 0.415, 0.425, 0.435, 0.445,
                0.455, 0.465, 0.475, 0.485, 0.495, 0.505, 0.515, 0.525, 0.535,
                0.545, 0.555, 0.565, 0.575, 0.585, 0.595, 0.605, 0.615, 0.625,
                0.635, 0.645, 0.655, 0.665, 0.675, 0.685, 0.695, 0.705, 0.715,
                0.725, 0.735, 0.745, 0.755, 0.765, 0.775, 0.785, 0.795, 0.805,
                0.815, 0.825, 0.835, 0.845, 0.855, 0.865, 0.875, 0.885, 0.895]),
 'l1_ratio': array([0.1 , 0.11, 0.12, 0.13, 0.14, 0.15, 0.16, 0.17, 0.18, 0.19,
                  0.2 ,
                  0.21, 0.22, 0.23, 0.24, 0.25, 0.26, 0.27, 0.28, 0.29, 0.3 , 0.31,
                  0.32, 0.33, 0.34, 0.35, 0.36, 0.37, 0.38, 0.39, 0.4 , 0.41, 0.42,
                  0.43, 0.44, 0.45, 0.46, 0.47, 0.48, 0.49])}
```

```
[30]: cv = RepeatedKFold(n_splits=5, n_repeats=2, random_state=42)
      print(cv)
```

```
RepeatedKFold(n_repeats=2, n_splits=5, random_state=42)
```

```
[31]: elasticNet = ElasticNet()
      Srch = GridSearchCV(elasticNet, parameters, scoring='neg_mean_squared_error',
      ↪cv=cv, n_jobs=-1)
      result = Srch.fit(train_set, y_train)
```

```
[32]: print(result)
      print(result.best_score_)
      print(result.best_params_)

      best_elastic = result.best_estimator_
      print(best_elastic)
```



```

GridSearchCV(cv=RepeatedKFold(n_repeats=2, n_splits=5, random_state=42),
            estimator=ElasticNet(), n_jobs=-1,
            param_grid={'alpha': array([0.005, 0.015, 0.025, 0.035, 0.045,
0.055, 0.065, 0.075, 0.085,
            0.095, 0.105, 0.115, 0.125, 0.135, 0.145, 0.155, 0.165, 0.175,
            0.185, 0.195, 0.205, 0.215, 0.225, 0.235, 0.245, 0.255, 0.265,
            0.275, 0.285, 0.295, 0.305, 0.315, 0.325, 0.335, 0.345, 0.355,
            0.36...
            0.725, 0.735, 0.745, 0.755, 0.765, 0.775, 0.785, 0.795, 0.805,
            0.815, 0.825, 0.835, 0.845, 0.855, 0.865, 0.875, 0.885, 0.895]),
            'l1_ratio': array([0.1 , 0.11, 0.12, 0.13, 0.14, 0.15,
0.16, 0.17, 0.18, 0.19, 0.2 ,
            0.21, 0.22, 0.23, 0.24, 0.25, 0.26, 0.27, 0.28, 0.29, 0.3 , 0.31,
            0.32, 0.33, 0.34, 0.35, 0.36, 0.37, 0.38, 0.39, 0.4 , 0.41, 0.42,
            0.43, 0.44, 0.45, 0.46, 0.47, 0.48, 0.49])}),
            scoring='neg_mean_squared_error')
-1047.2692562635616
{'alpha': 0.005, 'l1_ratio': 0.48999999999999977}
ElasticNet(alpha=0.005, l1_ratio=0.48999999999999977)

```

9 Comprobación de error cuadrático medio

Comprueba con la muestra de validación con cuál de los tres modelos se obtiene un menor error cuadrático medio.

9.1 Modelo 1

```

[33]: test_set = sm.add_constant(test_set, prepend=True)
      y_predict_model_1 = model1.predict(test_set)
      Mean_Sq_Error = mean_squared_error(y_predict_model_1, y_test)
      print(Mean_Sq_Error)

```

750.5664746002329

9.2 Modelo 2

```

[34]: test_set.drop(columns=['c2', 'c16', 'c17', 'c18', 'c19', 'c20', 'c21'], inplace=True)
      y_predict_model_2 = model2.predict(test_set)

```

```

[36]: Mean_Sq_Error_2 = mean_squared_error(y_predict_model_2, y_test)
      print(Mean_Sq_Error_2)

```

803.4930275424474

9.3 Modelo 3

```
[37]: y_predict_model_3 = best_elastic.predict(test_set)
```

```
[39]: Mean_Sq_Error_3 = mean_squared_error(y_predict_model_3, y_test)
print(Mean_Sq_Error_3)
```

796.5599792561268

10 Resultados de predicciones

```
[48]: resultados_dict = {
    'y_true' : y_test,
    'pred_modelo_1': y_predict_model_1,
    'pred_modelo_2': y_predict_model_2,
    'pred_modelo_3': y_predict_model_3,
}

resultados_df = pd.DataFrame(resultados_dict)

resultados_df
```

```
[48]:
```

	y_true	pred_modelo_1	pred_modelo_2	pred_modelo_3
132	333.020602	337.500686	337.793499	336.921428
148	48.171506	70.577423	68.063862	67.811020
93	-2.486599	6.473093	12.315322	12.395744
180	117.144811	98.533880	100.483649	100.065695
15	219.582261	179.860269	178.777864	178.464100
115	118.285486	93.841975	92.082986	91.826115
172	35.675697	92.450567	99.545055	99.347175
209	-165.223912	-159.573230	-164.636243	-164.098968
75	-102.658248	-111.493527	-112.536782	-112.142209
142	78.210423	44.007023	51.301069	51.217902
100	136.480734	87.929544	86.054299	85.798261
30	239.090672	248.634253	246.631635	245.976291
190	-41.514453	-31.521500	-20.886600	-20.791891
9	-395.573814	-401.397174	-395.046826	-394.028853
67	31.789654	39.031722	38.286235	38.367161
218	373.228970	407.431108	400.094889	398.936383
175	-336.495524	-362.971016	-365.631178	-364.613181
18	-99.469821	-152.123828	-155.469751	-155.049234
197	-139.163819	-137.783906	-149.326964	-149.080743
66	-107.825197	-118.043561	-112.738986	-112.551307

Según los resultados mostrados, el mejor modelo fué el Modelo 1, ya que tiene el valor de mean squared error más bajo, luego dió mejor resultado el modelo 3, teniendo en consideración el uso de elastic net. Por último el modelo que tenía mas alto el valor de mean squared error es del

modelo 2. Posiblemente estos resultados pueden mejorar teniendo una mayor cantidad de datos, ya que la cantidad de datos utilizados en realidad han sido pocos.

References

- [1] Aurélien Géron. *Hands-on Machine Learning with Scikit-Learn and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems*. O'Reilly Media.
- [2] Scikit-Learn. Gridsearchcv. https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.GridSearchCV.html.
- [3] Scikit-Learn. Repeatedkfold. https://sklearn.org/modules/generated/sklearn.model_selection.RepeatedKFold.html.